

# 1 542 AIDevSet Project

2  
3 YUE WANG\*, University of British Columbia at Okanagan, Canada  
4 YUSEN RONG†, University of British Columbia at Okanagan, Canada

5  
6 **ACM Reference Format:**

7 Yue Wang and Yusen Rong. 2025. 542 AIDevSet Project. 1, 1 (December 2025), 7 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

8  
9 **1 Introduction**

10 With the rise of artificial intelligence (AI) in recent years, there has been a lot of discussion about  
11 if and how soon AI can replace software engineers in most companies. Historically, the code  
12 generated by AI can be questionable, especially if the language is relatively new or not as popular.  
13 Several analysis will need to be done on the capabilities of AI agents before even broader adoption  
14 by companies, and is probably a yearly exercise for every company to evaluate effectiveness and  
15 performance. In this report, we will use some of the AIDev datasets to analyze some performance  
16 and adoption metrics to see where the state of AI agents is currently. The GitHub project for this  
17 report is located at <https://github.com/yuewangse/542Project>.

18  
19 **2 Analyzing of AI Agent Usage VS Coding Language**

20 The first question that comes to mind is which coding language uses AI agent Pull Requests more  
21 often? Is there a specific tendency to adopt AI agents more in some languages vs others? This data  
22 is very interesting and might give us some insight on what language AI agents work best, where  
23 they need improvement, and if companies that utilize a certain language should start looking at  
24 adopting AI agent usage if they have not already.

25  
26 **2.1 Methodology**

27 To determine the frequency of AI agent usage per language, we will mainly use two of the AIDev  
28 datasets, pull\_requests and repository. We first checked that there are no Na values inside each of  
29 the repositories so that we can combine the two based on the repo\_id column, unfortunately we  
30 found that around 37 out of 2807 repositories do not have data on what language they use, and with  
31 such a low number affected and no reliable way to collect that data ourselves, we simply discarded  
32 those rows of data. After the merge we confirmed that 183 of the 33596 rows of pull request data  
33 was also lost, which is a small number as well and was acceptable to us. With the cleaned data,  
34 we can then calculate and plot the frequencies of repos per language, the frequencies of prs per

---

35  
36 \*General Report structure, wrote Introduction, Conclusion, section 2, overviews for section 3 and 4, wrote code for section 2  
37 and subsection, proof reader

38 †Wrote methodology and analysis for section 3 and 4, wrote code for section 3 and 4 and subsections

---

39 Authors' Contact Information: Yue Wang, [yuewang@student.ubc.ca](mailto:yuewang@student.ubc.ca), University of British Columbia at Okanagan, Kelowna,  
40 British Columbia, Canada; Yusen Rong, [yusenrong46@gmail.com](mailto:yusenrong46@gmail.com), University of British Columbia at Okanagan, Kelowna,  
41 British Columbia, Canada.

---

42 Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee  
43 provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the  
44 full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored.  
45 Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires  
46 prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

47 © 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

48 ACM XXXX-XXXX/2025/12-ART

49 <https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

language, and finally divide the two statistics to get average number of prs per repo per language to estimate language adoption.

## 2.2 Results

The generated frequency graphs are as follows, showing only top 20 frequencies to keep it legible.




Fig. 1. Frequency Graph Depicting Number of Repos by Language




Fig. 2. Frequency Graph Depicting Number of AI Agent Prs by Language




Fig. 3. Frequency Graph Depicting Average Number of AI Agent Prs per Repo by Language

### 99    2.3 Analysis

100 Looking at Figs. 1 and 2 above, we can see that the most popular languages to use in code repositories  
 101 do not match the most popular languages to use AI agent prs on, which could have several causes.  
 102 One of our guesses is that some languages are easier for ai agents to work with and thus have  
 103 better models and be more developer friendly, while others are the opposite. Another guess is that  
 104 companies that use those languages are more optimistic about AI usage and have pushed their  
 105 employees to integrate more with AI agents than others.

106 The final Fig. 3 above shows the average amount of AI agent pull requests per repo separated  
 107 by language. This can serve as an estimate of AI agent adoption per language for consideration,  
 108 and might help a company decide whether they should follow suit/continue on their path of AI  
 109 integration if they also use the same language. Do note however that this graph can be misleading,  
 110 as the language with the largest adoption of AI agents, Circom, is used in an extremely small  
 111 number of repos, so it might just be the effect one person who is really into AI adoption. We  
 112 would recommend only consulting the results for the more popular programming languages and  
 113 not all of them.

## 115    3 Analyzing Change Complexity of Different AI Agents

116 In this question, we will study how various factors of change complexity of different AI agents  
 117 affect the time to merge and acceptance rate of AI agent created pull requests. This can be useful to  
 118 know what factors influence acceptance rate and merge times to help a company decide which  
 119 AI agent they might want to use or migrate to best fit the company's goals.

### 121    3.1 Methodology

122 For this question, we study how the size and complexity of AI-generated pull requests relate to  
 123 their probability of being merged. We work with the AIDev dataset at the pull-request level and  
 124 focus only on PRs created by AI agents (OpenAI Codex, Copilot, Devin, Cursor, Claude Code).

125 We first construct a base pull-request table from `pull_request.parquet` by renaming the  
 126 primary key to `pr_id` and the agent column to `agent_name`. We then define a binary outcome  
 127 variable `merged_flag`, which is set to 1 if the PR has a non-missing `merged_at` timestamp and 0  
 128 otherwise. This serves as our measure of "acceptance": a PR is considered accepted if it is merged  
 129 into the main branch.

130 Next, we build change-complexity features using the `pr_commit_details.parquet` table. For  
 131 each `pr_id`, we aggregate file-level statistics to obtain:

- 132    • `total_additions`: total lines added across all files,
- 133    • `total_deletions`: total lines removed,
- 134    • `files_changed`: number of unique files touched,
- 135    • `total_changes`: total lines changed (sum of changes across files).

136 We then merge these complexity features back into the PR table. To reduce skew, we define  
 137  $\log_{\text{total\_changes}} = \log(1 + \text{total\_changes})$ . Missing counts are treated as zeros, which corre-  
 138 sponds to PRs with very small or trivial diffs.

139 As an initial descriptive step, we compute per-agent summaries: number of PRs, mean and  
 140 median `total_changes`, and the empirical merge rate (`merged_flag` mean). This shows how each  
 141 agent behaves in terms of typical patch size and acceptance rate.

142 To formally quantify the relationship between complexity, agent identity, and merge outcomes,  
 143 we fit a logistic regression model. The response is `merged_flag` (1 = merged, 0 = not merged). The  
 144 predictors are:

- 145    • numeric features: `log_total_changes`, `files_changed`;

- 148     • categorical feature: `agent_name`, encoded via one-hot encoding.

149 We standardize numeric features with `StandardScaler` and encode `agent_name` with `OneHotEncoder`  
 150 inside a single scikit-learn pipeline. The dataset is split into training and test sets using an 80/20  
 151 stratified split to preserve the merge/non-merge ratio. To evaluate model stability and out-of-sample  
 152 performance, we use 10-fold stratified cross-validation and report the mean ROC AUC and its  
 153 standard deviation, as well as the ROC AUC on the held-out test set. Finally, we extract the logistic  
 154 regression coefficients and convert them into odds ratios to interpret the effect of complexity and  
 155 agent identity on merge probability.  
 156

### 157     3.2 Result

158 Descriptively, the five AI agents show clear differences in both change size and acceptance rates.  
 159 OpenAI Codex submits the largest number of PRs and tends to produce relatively small, focused  
 160 changes: its median `total_changes` is on the order of tens of lines, and its merge rate is high  
 161 (approximately 83%). In contrast, Copilot and Devin submit substantially larger changes on average,  
 162 with higher median `total_changes`, but achieve lower merge rates (around 43% for Copilot and  
 163 54% for Devin). Cursor and Claude Code occupy a middle ground: they typically propose medium-  
 164 to-large patches with merge rates between those of Codex and Copilot.  
 165

166 The logistic regression model that combines change complexity (`log_total_changes`, `files_changed`)  
 167 and agent identity (`agent_name`) achieves a mean ROC AUC of about 0.716 (standard deviation  
 168 roughly 0.004) under 10-fold stratified cross-validation. On a held-out 20% test set, the ROC AUC is  
 169 approximately 0.717. This means that, given a randomly chosen merged PR and a non-merged PR,  
 170 the model gives a higher predicted merge probability to the merged PR about 71–72% of the time.  
 171 Overall accuracy on the test set is around 71%, with the model performing better for the merged  
 172 class (precision ≈ 0.83, recall ≈ 0.75) than for the non-merged class.

173 Looking at the coefficients, `log_total_changes` has a negative effect on merge odds: its estimated  
 174 odds ratio is roughly 0.8, meaning that a one-standard-deviation increase in log total changes  
 175 reduces the odds of being merged by about 20%, holding agent identity and other factors fixed. The  
 176 coefficient on `files_changed` is small and slightly positive (odds ratio close to 1.05), suggesting  
 177 that once we know how many lines are changed, the number of files touched adds relatively little  
 178 additional information.

179 Agent identity also remains important after controlling for change complexity. The indicator  
 180 associated with OpenAI Codex has a strong positive coefficient, with an odds ratio of about 2.5,  
 181 implying that Codex PRs have more than double the merge odds of an “average” encoded agent  
 182 at the same complexity level. Cursor has a mild positive effect (odds ratio slightly above 1), while  
 183 Claude Code is roughly neutral (odds ratio just below 1). By contrast, Devin and especially Copilot  
 184 have negative coefficients, with odds ratios of about 0.66 and 0.41 respectively, indicating that their  
 185 PRs are less likely to be merged than similarly complex Codex or Cursor PRs.

### 186     3.3 Analysis

187 The results suggest two main findings for change complexity and AI agents.

188 First, change complexity clearly matters. Larger, more complex PRs—as measured by `total_changes`—  
 189 are less likely to be merged, even after controlling for which agent created them. The negative odds  
 190 ratio for `log_total_changes` formalizes a pattern that many practitioners intuitively expect: big  
 191 diffs are harder to review, more likely to break existing behaviour, and therefore more likely to be  
 192 rejected or left unmerged. The very small effect of `files_changed` once total lines changed are  
 193 known suggests that reviewers care primarily about the overall amount of change rather than the  
 194 exact number of files it is spread across.

197 Second, agent identity remains an important predictor of merge success over and above complexity. OpenAI Codex PRs are both smaller on median and substantially more likely to be merged,  
 198 while Copilot and Devin PRs are larger and have noticeably lower merge odds. Cursor and Claude  
 199 Code sit in between. Since the model controls for change size, these differences cannot be explained  
 200 purely by the fact that some agents produce larger patches. Instead, they indicate that maintainers  
 201 respond differently to different agents' outputs, possibly because of perceived code quality, style,  
 202 or stability of the changes.  
 203

204 From a practical perspective, this analysis provides a simple way for an engineering organization  
 205 to think about adopting or migrating between AI coding agents. If a team cares primarily about  
 206 maximizing the fraction of AI-generated PRs that get merged with minimal human friction, then  
 207 an agent that tends to produce smaller, more focused edits with higher conditional merge odds  
 208 (such as Codex in our dataset) appears more attractive. If a team instead wants aggressive refactors  
 209 or large-scale changes, they should expect lower acceptance rates and may need to invest more in  
 210 code review and testing workflows.

211 Finally, we note that this question focuses on merge probability as our main outcome. Time-  
 212 to-merge is also an important part of developer productivity, but it depends strongly on human  
 213 review dynamics. In the next research question, we explicitly model review intensity (number of  
 214 reviewers, timing of first review, and review activity) to see how human oversight interacts with  
 215 agent identity and change complexity to shape both merge decisions and review effort.

## 216 4 Analyzing Review Intensity of Different AI Agents

217 For the final question, we will analyze how human scrutiny changes based on which AI agent  
 218 constructed the pull request. We can measure the number of reviewers, total review comments,  
 219 and time to first review for each AI agent to assess their performance compared to each other. This  
 220 can also be helpful to companies deciding which AI agent they should work with.

### 221 4.1 Methodology

222 For RQ3, we study how human oversight signals (reviews and comments) relate to pull request  
 223 (PR) merge outcomes for AI-generated PRs. The unit of analysis is one PR, meaning each row in  
 224 the final dataset corresponds to a single PR.

225 We construct a PR-level table from `pull_requests.parquet` and define the outcome variable  
 226 `merged_flag`, where a PR is labeled 1 if `merged_at` is present and 0 otherwise. We keep only PRs  
 227 with a finalized state (closed PRs) to ensure every PR has a resolved outcome.

228 Next, we build PR-level complexity features using `pr_commit_details.parquet`, aggregated  
 229 by `pr_id`. These include `total_changes` and `files_changed`. Because PR size is highly skewed,  
 230 we define `log_total_changes = log(1 + total_changes)`.

231 To represent oversight, we aggregate multiple human-interaction tables at the PR level. From  
 232 `pr_reviews.parquet`, we compute counts such as `n_reviews`, `n_reviewers`, `n_approvals`, and  
 233 `n_changes_requested`. From `pr_comments.parquet`, we compute `n_pr_comments`. From `pr_review_comments`  
 234 we compute `n_inline_review_comments`. We also compute `time_to_first_review_hours` us-  
 235 ing the PR creation time and the timestamp of the first review. If a PR has no review, we set  
 236 `time_to_first_review_hours` to -1 and include a binary indicator `has_review`.

237 We then train a Random Forest classifier to predict `merged_flag` using both complexity and  
 238 oversight features, plus the categorical feature `agent_name` (one-hot encoded). Missing numeric  
 239 values are imputed using the median within a single scikit-learn pipeline. We use an 80/20 stratified  
 240 train/test split and evaluate with 10-fold stratified cross-validation on the training set. Performance  
 241 is reported using ROC AUC, and we also report the confusion matrix and classification metrics on  
 242 the held-out test set.

## 246 4.2 Result

247 The model uses 25,014 PRs for training and 6,254 PRs for testing. Under 10-fold cross-validation,  
 248 the mean ROC AUC is approximately 0.751 (standard deviation about 0.011). On the held-out test  
 249 set, the ROC AUC is 0.741 and the overall accuracy is 0.79.  
 250

Metric	Value
Train size	25,014
Test size	6,254
10-fold CV ROC AUC (mean $\pm$ sd)	$0.751 \pm 0.011$
Test ROC AUC	0.741
Test accuracy	0.79


251 252 253 254 255 256 257 258 259 Table 1. Random Forest performance for predicting merge outcome.  
 260  
 261

262 The confusion matrix (rows=true, cols=pred) is shown in Table 2. The model predicts merged  
 263 PRs well (class 1 recall 0.88) but is weaker at identifying non-merged PRs (class 0 recall 0.48).  
 264

	Pred 0	Pred 1
True 0	700	751
True 1	574	4229

265 266 267 268 269 270 271 272 Table 2. Confusion matrix on the test set (rows=true, cols=pred).

273 Figure 4 shows merge rate versus an oversight intensity measure (binned). Figure 5 shows  
 274 the most important features in the Random Forest model. The top feature importances include  
 275 log\_total\_changes (0.4165), files\_changed (0.1347), n\_pr\_comments (0.0964), agent\_name\_OpenAI\_Codex  
 276 (0.0805), n\_approvals (0.0789), and time\_to\_first\_review\_hours (0.0548).  
 277



292 293 Fig. 4. Merge rate versus oversight intensity (binned) from the RQ3 notebook.  
 294




Fig. 5. Top Random Forest feature importances from the RQ3 notebook.

### 4.3 Analysis

Two patterns stand out.

First, PR complexity dominates the prediction problem. The single most important feature is `log_total_changes`, and `files_changed` is also highly ranked. This is consistent with a simple workflow reality: larger PRs are harder to review, more likely to introduce risk, and therefore less likely to be merged quickly or at all.

Second, oversight signals still matter. Features like `n_pr_comments`, `n_approvals`, and `time_to_first_review` carry meaningful importance, and Figure 4 suggests merge rate changes across different levels of oversight intensity. However, this relationship should be interpreted as correlational rather than causal. Oversight is often a response to underlying PR quality, complexity, urgency, or team norms, so high oversight does not necessarily *cause* a merge. In addition, the test-set confusion matrix shows the model is much better at detecting merged PRs than non-merged PRs, which is consistent with class imbalance where merged PRs are the majority.

## 5 Conclusion

We have successfully analyzed several aspects of AI agents and how they compare to one another, as well as the usage statistics of the agents based on language. The results of the analysis in this report should help a company decide if it is time to start adopting AI agents in their workflow, and if so which ones to consider first.